

Data Science Practicals: Final Assignment

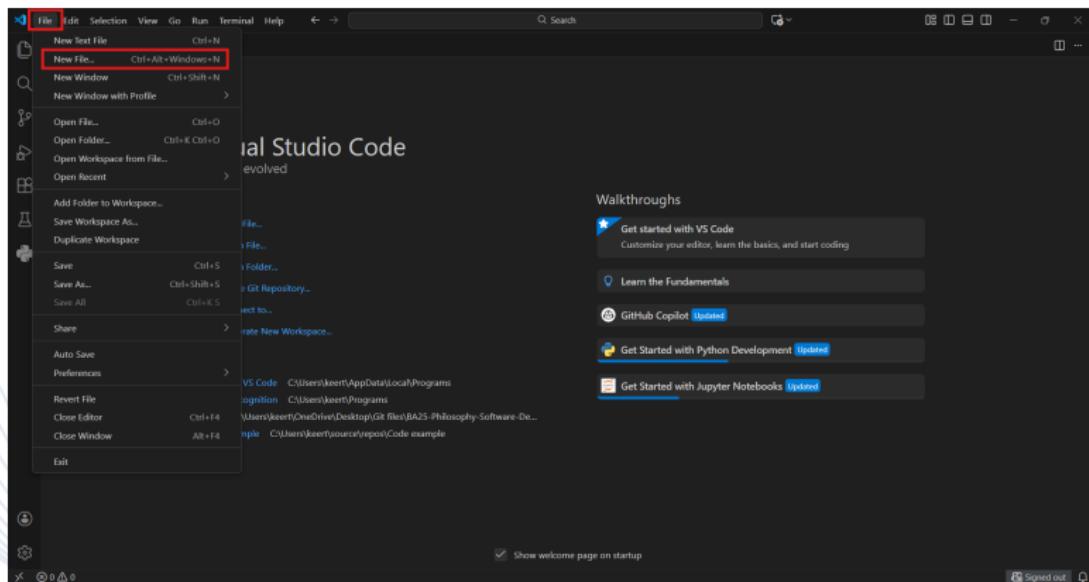
Adil Ibraheem Koyava, Dieter Kalwa & Keerti Belmane

16. Dezember 2025

- 1 Creation of .ipynb file
- 2 Final Assignment: Overview
- 3 Task 1: Describe the Dataset
- 4 Task 2: Predict Demand from Conditions
- 5 Task 3: 30-Day Forecast
- 6 Final Assignment: Evaluation

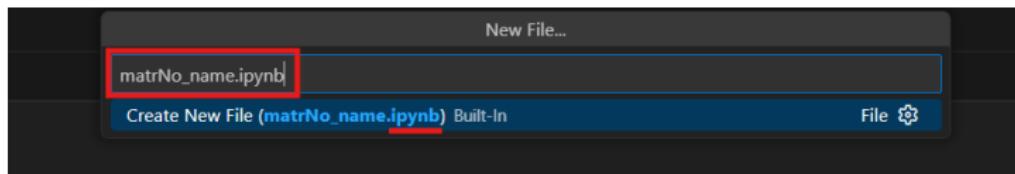
Creation of .ipynb file

Creating the file



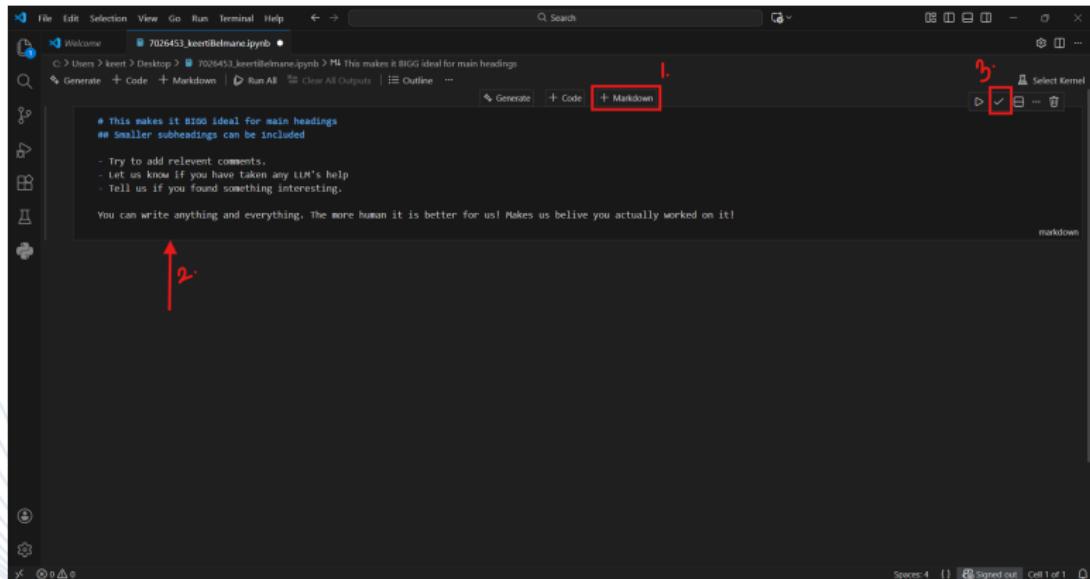
- Click on "File"
- Click on "New File..."

Naming the file



- Name the file with your **MatriculationNumber_Name.ipynb**.
- **Not .py!!** Save it as **.ipynb**, this makes it the interactive jupyter notebook that we are looking for.

Adding Markdown Comments



The screenshot shows a Jupyter Notebook interface with a single Markdown cell containing the following text:

```
# This makes it BIGG ideal for main headings
## Smaller subheadings can be included

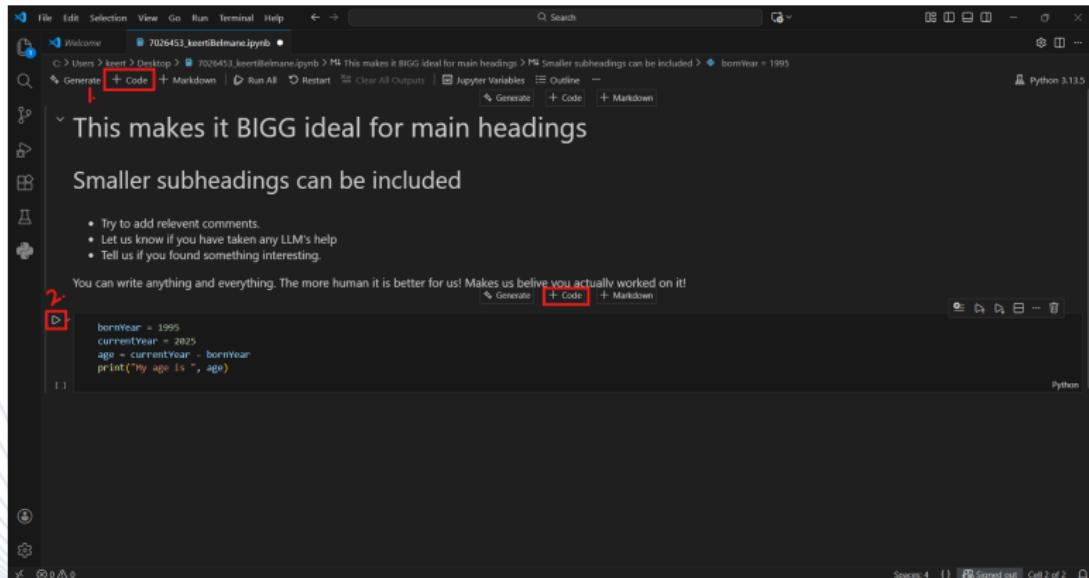
- Try to add relevant comments.
- Let us know if you have taken any LHM's help
- Tell us if you found something interesting.

You can write anything and everything. The more human it is better for us! Makes us believe you actually worked on it!
```

A red arrow points upwards from the bottom of the slide towards the save icon in the top right corner of the notebook window. A red box highlights the "Markdown" button in the toolbar above the cell.

- Click on the Markdown, and add your comments.
- Click on the tick mark on the right up corner of the markdown cell. This will save it.

Adding Code cells



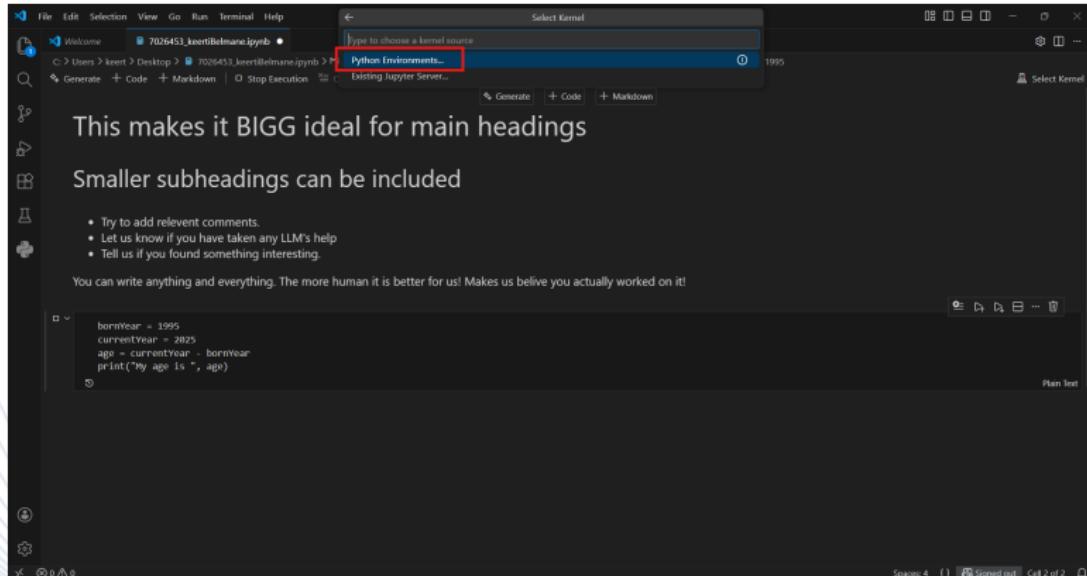
The screenshot shows a Jupyter Notebook interface with the following details:

- Header:** File, Edit, Selection, View, Go, Run, Terminal, Help, Search.
- Toolbar:** Welcome, 7026453_keertBelmane.ipynb, Generate, + Code (highlighted with a red box), + Markdown, Run All, Restart, Clear All Outputs, Jupyter Variables, Outline, Python 3.13.5.
- Content Area:**
 - A heading: "This makes it BIGG ideal for main headings".
 - A subheading: "Smaller subheadings can be included".
 - A bulleted list:
 - Try to add relevant comments.
 - Let us know if you have taken any LLM's help.
 - Tell us if you found something interesting.
 - A message: "You can write anything and everything. The more human it is better for us! Makes us believe you actually worked on it!"
 - A code cell with Python code:

```
bornYear = 1995
currentYear = 2025
age = currentYear - bornYear
print("My age is ", age)
```
- Bottom Status Bar:** Spaces: 4, Signed out, Cell 2 of 2.

- Click on the code, to create the code cell. Here write all the code you want to implement.
- Click on the execute button, to execute the code.

Setting the python variable



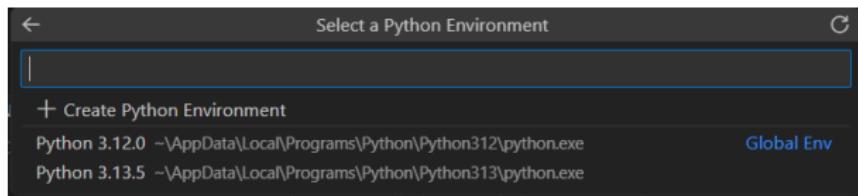
The screenshot shows a Jupyter Notebook interface. At the top, a 'Select Kernel' dialog box is open, with 'Python Environments...' highlighted. The main notebook area displays the following code:

```
bornYear = 1995
currentYear = 2025
age = currentYear - bornYear
print("My age is ", age)
```

Below the code, there is a note: "You can write anything and everything. The more human it is better for us! Makes us believe you actually worked on it!"

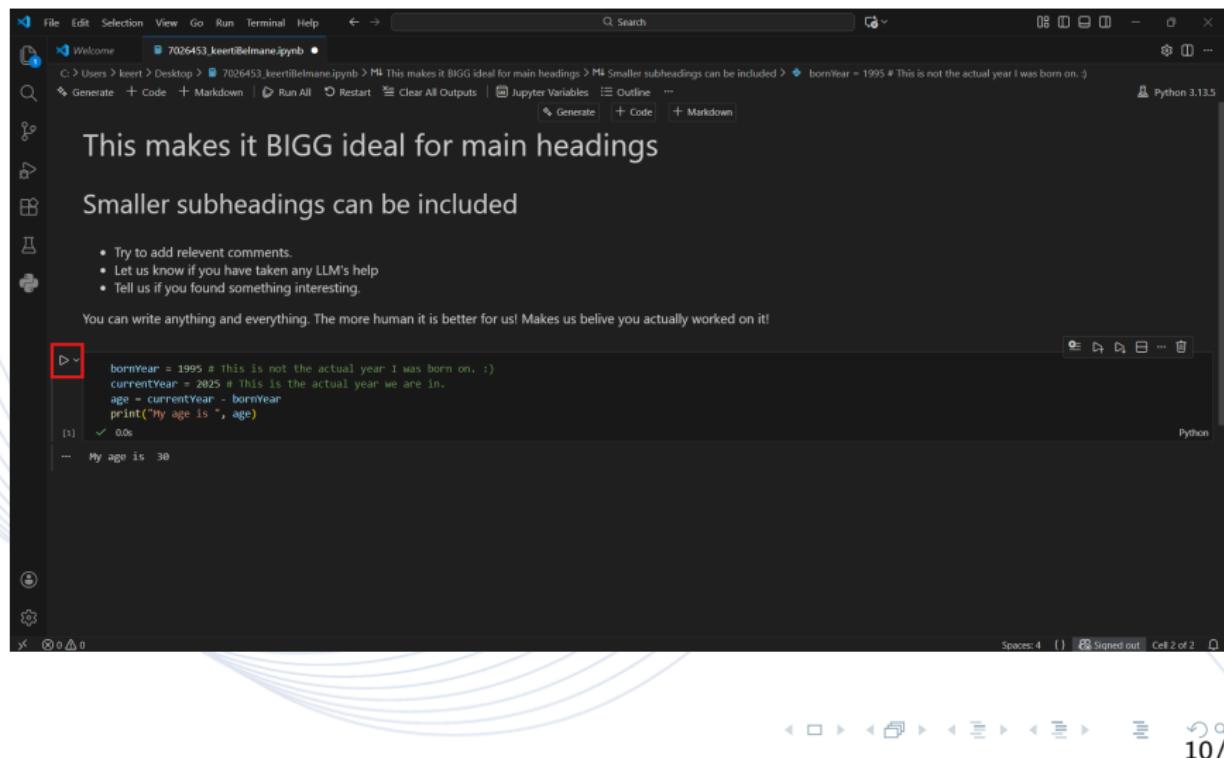
- 1st time when you try to execute, you will have to set your python environment.
- Click on Python Environment

Select the python version



- Select the python version you would like to use.

Happy Coding to you! :) All the best



The screenshot shows a Jupyter Notebook interface with the following details:

- File Bar:** File, Edit, Selection, View, Go, Run, Terminal, Help.
- Toolbar:** Welcome, 7026453_keertilelmane.ipynb, Search, Python 3.13.5.
- Code Cell:** C:\> Users\> keert > Desktop > 7026453_keertilelmane.ipynb > M4 This makes it BIGG ideal for main headings > M4 Smaller subheadings can be included > bornYear = 1995 # This is not the actual year I was born on. ;)
- Cell Buttons:** Generate, + Code, + Markdown, Run All, Restart, Clear All Outputs, Jupyter Variables, Outline, % Generate, + Code, + Markdown.
- Section Headers:**

This makes it BIGG ideal for main headings

Smaller subheadings can be included
- List Item:** Try to add relevant comments.
Let us know if you have taken any LLM's help
Tell us if you found something interesting.
- Text:** You can write anything and everything. The more human it is better for us! Makes us believe you actually worked on it!
- Code Output:** bornYear = 1995 # This is not the actual year I was born on. ;)
currentYear = 2025 # This is the actual year we are in.
age = currentYear - bornYear
print("My age is ", age)
[1] ✓ 0.0s
-- My age is 30
- Bottom Status:** Spaces: 4, Signed out, Cell 2 of 2, Page Number: 10/37.

Final Assignment: Overview

Bike Sharing Demand Dataset I

The final project uses a real-world bike sharing demand dataset (CSV file) based on historical usage of a public bike rental system.

The dataset contains:

- A **target variable** count: number of rented bikes in a given time period.
- **Time information**: a datetime column (date and hour or date and day).
- **Context features**: weather, temperature, humidity, windspeed, season, holiday and working day indicators.

The overall goal is to build:

- A model that predicts bike demand from given conditions (“what-if” prediction).
- A time-series model that forecasts the next 30 days of total demand, following the pipeline from the lecture.

Dataset Variables I

Variable Name	Type	Description
datetime	Datetime	Date and time of the observation (hourly or daily).
count	Numeric	Total number of bikes rented (target variable).
season	Categorical	Season of the year (1=spring, 2=summer, 3=fall, 4=winter).
month	Numeric	Month of the year (1–12).
hour	Numeric	Hour of the day (0–23), if hourly data.
weekday	Categorical	Day of the week (0=Monday, ..., 6=Sunday).
holiday	Binary	Whether the day is a holiday (0=no, 1=yes).

Tabelle: Bike Sharing Dataset Variables (1/2)

Dataset Variables II

Variable Name	Type	Description
workingday	Binary	Whether the day is a working day (0=no, 1=yes).
weather	Categorical	Weather condition code (1=clear, 2=mist, 3=light rain/snow, 4=heavy rain/snow).
temp	Numeric	Temperature in °C.
atemp	Numeric	“Feels like” temperature in °C.
humidity	Numeric	Relative humidity (%).
windspeed	Numeric	Wind speed (km/h or similar unit).
casual	Numeric	Number of casual users. <i>Do not use in Task 2 or 3.</i>
registered	Numeric	Number of registered users. <i>Do not use in Task 2 or 3.</i>

Tabelle: Bike Sharing Dataset Variables (2/2)

Dataset Variables III

Note: The `casual` and `registered` columns are derived from `count` (they sum to `count`) and should **never** be used as input features to avoid data leakage.

Task 1: Describe the Dataset

Task 1: Data Understanding I

Goal: Understand the structure and basic properties of the dataset using Python (NumPy, pandas, Matplotlib/Seaborn).

What to do:

- ① Read the CSV file with `pandas.read_csv` and parse the `datetime` column.
- ② Report:
 - Number of rows and columns.
 - Time range covered by the data.
 - Target variable and list of feature variables (names and data types).
- ③ Create a **variable description table** (see above for reference).
- ④ Check for:
 - Missing values per column.
 - Duplicated rows (if any).

Task 1: Descriptive Statistics and Visualisation I

Descriptive statistics:

- For numeric variables: calculate mean, standard deviation, minimum, maximum, and quartiles (e.g. with `df.describe()`).
- For categorical variables: show frequency tables or bar charts (e.g. distribution of seasons or weather types).
- Check for missing values in each column and comment on how you will handle them.

Visualisation:

- Plot the time series of total bike demand (`count`) over the full period.
- Plot distributions of key numeric variables (e.g. histograms of `temp`, `humidity`, `windspeed`).
- Plot aggregated demand by season, day of week or hour of day (e.g. bar charts or line plots).

Task 1: Descriptive Statistics and Visualisation II

Write a short text summary (3–5 sentences) describing main patterns you observe (seasonality, daily patterns, influence of weather, etc.).

Task 2: Predict Demand from Conditions

Task 2: Supervised Regression Setup II

Goal: Build a supervised regression model that predicts bike demand count from given conditions (features such as weather, temperature, time of day).

Input and target:

- Target variable: count (total number of bikes rented in that period).
- Example input features:
 - Calendar: season, month, day of week, hour of day, holiday, working day.
 - Weather: weather situation code, temperature, humidity, windspeed.
- Do **not** use columns that directly leak the target (e.g. casual, registered if present).

Task 2: Supervised Regression Setup II |

Preprocessing:

- Encode categorical variables (e.g. one-hot encoding for season, weather, weekday).
- Split the data into training, validation and test sets (for example 70 % / 10 % / 20 %) and scale numeric features using statistics from the training set only.

Task 2: Train Model and Implement User Prediction I

Model training:

- Choose at least one regression model using a machine learning library such as scikit-learn, TensorFlow/Keras or PyTorch.
- Examples: Linear Regression, Random Forest, Gradient Boosting, XGBoost, Multi-Layer Perceptron (dense neural network).
- Train the model on the training set and tune hyperparameters / monitor performance with the validation set.

Evaluation:

- Evaluate the final model on the test set using at least:
 - RMSE (Root Mean Squared Error).
 - MAE (Mean Absolute Error).
 - R^2 (coefficient of determination).
- Plot predicted vs actual demand on the test set (scatter plot and/or line plot).

Task 2: Train Model and Implement User Prediction II

User-input prediction function:

- Implement a Python function (or cell) that:
 - Accepts user inputs (e.g. season, weather, temp, humidity, windspeed, hour, weekday, `is_holiday`, `is_workingday`).
 - Applies the same preprocessing and scaling steps.
 - Uses the trained model to output a predicted count of bikes.
- Demonstrate several example predictions for different scenarios and briefly interpret the results.

Task 3: 30-Day Forecast

Task 3: Time-Series Pipeline I

Goal: Build a time-series model that forecasts bike demand for the next 30 days (or 30 time steps), following the forecasting pipeline from the lecture.

Data preparation:

- Use the chronological order of the data to create a time-based split:
 - Training period (e.g. first 70 % of the timeline).
 - Validation period (e.g. next 15 %).
 - Test period (e.g. last 15 %).
- Engineer **time-series features**:
 - Lag features of count (e.g. 1, 7, 24, 168 steps back).
 - Rolling statistics (e.g. 7-step and 30-step rolling mean/standard deviation of count) using shifted values to avoid data leakage.
- Scale features using statistics from the training period only.

Task 3: Forecasting Model and Metrics II

Modeling options:

- Sequential models (e.g. LSTM, GRU, 1D CNN, sequence-to-one dense networks):
 - Create input windows of length 30: past 30 time steps as input, next step as target.
 - Input shape: (batch, 30, num_features).
- Tabular models (e.g. Random Forest, Gradient Boosting, XGBoost, dense neural nets):
 - Use lag and rolling features directly as a 2D feature matrix.

Task 3: Forecasting Model and Metrics II I

Training and evaluation:

- Train the model on the training period and monitor hyperparameters / early stopping with the validation period.
- Evaluate performance on the test period using at least:
 - RMSE (Root Mean Squared Error).
 - MAE (Mean Absolute Error).
 - MAPE (Mean Absolute Percentage Error).
- Plot true vs predicted demand on the test period as a time series and compare to a naive baseline (e.g. last observed value or moving average forecast).

Task 3: 30-Day Forecast and Interpretation II

Multi-step forecast (30 days):

- Starting from the end of the known data, generate predictions step by step:
 - At each step, use the most recent history (and features) as model input.
 - Append the predicted demand to the series and recompute necessary lag/rolling features.
 - Repeat until a 30-day (or 30-step) forecast horizon is reached.
- Plot the full history of demand together with the 30-day forecast on a single chart.

Task 3: 30-Day Forecast and Interpretation II |

Interpretation and comparison:

- Discuss whether the forecast captures observed patterns (e.g. seasonality, weekday/weekend effects, weather influence).
- Compare the forecasting model to:
 - The naive baseline.
 - The user-input regression model from Task 2 (when and why each is useful).
- Summarise the strengths and limitations of your approach in a short written conclusion.

Final Assignment: Evaluation

Declaration of Authorship I

I hereby declare that I, the undersigned, am the sole author of this submission. All sources consulted have been listed; all quotations and references have been properly cited. No version of this submission (in whole or in part) has been used previously for an academic degree or other examination.

I understand that false statements in this declaration may be punishable by law.

If ChatGPT or other LLM-based tools were used, clearly specify in the notebook:

- Which tool(s) were used (e.g. ChatGPT, Copilot)
- For what purpose (e.g. syntax help, explanation, debugging)
- Which sections or cells were influenced

Evaluation Criteria — Mark Allocation (100 points) I

Criterion	Pts
Authorship & LLM disclosure	4
Reproducibility	4
Error handling	7
Feature engineering	4
Data splitting	11
Data leakage prevention	7
Model architecture	7
Hyperparameters	11
Model evaluation	7
Final prediction	2
Comments and explanations of code	7
Code quality	7
Environment specification	7
Submission format (.ipynb)	7
Delighter	8
Total	100

Evaluation Criteria — Explanation (I) I

Criterion	Explanation
Authorship & LLM disclosure	A declaration of authorship is included. Any use of ChatGPT or similar tools is explicitly stated and described.
Reproducibility	The notebook executes top-to-bottom on a clean environment and produces consistent results. Random seeds are fixed where appropriate.
Error handling	Potential runtime errors are anticipated and handled (e.g. using try/except), with informative messages for the user.
Feature engineering	Meaningful features are added to the raw data (e.g. Fourier features for seasonality) and their purpose is explained.
Data splitting	Training, validation, and test sets are constructed correctly and suit the nature of the data (time-aware where applicable).
Data leakage prevention	Scaling, encoding, and feature transformations are fitted on training data only; validation and test data are never used improperly.

Evaluation Criteria — Explanation (II) I

Criterion	Explanation
Model architecture	A suitable TensorFlow/Keras/Other model is built with appropriate layers, activations, and input dimensions.
Hyperparameters	Key hyperparameters (learning rate, epochs, etc.) are identified and briefly motivated.
Model evaluation	Appropriate accuracy metrics (e.g. MAE, RMSE) are computed correctly and interpreted in plain language.
Final prediction	The model produces a clear final prediction in the required format (e.g. future values).
Comments and explanations of code	Markdown cells explain the workflow clearly, and code contains concise, helpful inline comments.

Evaluation Criteria — Explanation (III) I

Criterion	Explanation
Code quality	Variables are clearly named, imports are minimal and relevant, and Python conventions are followed.
Environment specification	The Python version and key library versions (e.g. NumPy, TensorFlow) are explicitly stated.
Submission format	The final submission is a valid .ipynb file that opens and runs without errors.
Delighter	A thoughtful enhancement beyond the minimum requirements (e.g. baseline comparison, diagnostics, visualisations).

Thank you
for your attention